

Project Report – ECE 564 Fall 2025

1024×1024 CNN Pipeline with 4×4 Convolution, Leaky-ReLU, and 2×2 Average Pooling

Author: Shivanesh Gunasekaran

Unity ID: sgunase2

Clock Period: 3.8 ns

1. Introduction

This project implements a complete hardware pipeline for a simplified Convolutional Neural Network (CNN) using SystemVerilog RTL. The design reads a 1024×1024 signed 8-bit input image from DRAM, performs:

1. **4×4 convolution**
2. **Leaky-ReLU activation**
3. **2×2 average pooling**

and finally writes the processed output back to DRAM.

The entire system follows the exact timing and memory interface protocol provided in the class project specification and uses a **pipelined architecture** capable of sustaining computation at the DRAM read throughput.

The design uses the optional **dual-port SRAM** as an internal scratchpad to buffer input rows and feed the convolution unit without stalling.

2. High-Level System Overview

The design consists of four major subsystems:

1. DRAM Input Loader FSM

Fetches kernel (16 bytes) and input pixels ($1024 \times 1024 \times 8$ -bit) from DRAM following burst-read protocol.

2. SRAM Data Re-organization Engine

Stores the input matrix into SRAM row buffers to enable fast sliding-window extraction for convolution.

3. Convolution + ReLU + Pooling Engine (Main Datapath)

For every valid 4×4 window:

- Performs 16 parallel multiplications
- Performs 15 additions across three pipeline stages
- Applies Leaky-ReLU
- Performs 2×2 average pooling (with padding logic)

Implemented as a **deeply pipelined compute engine**.

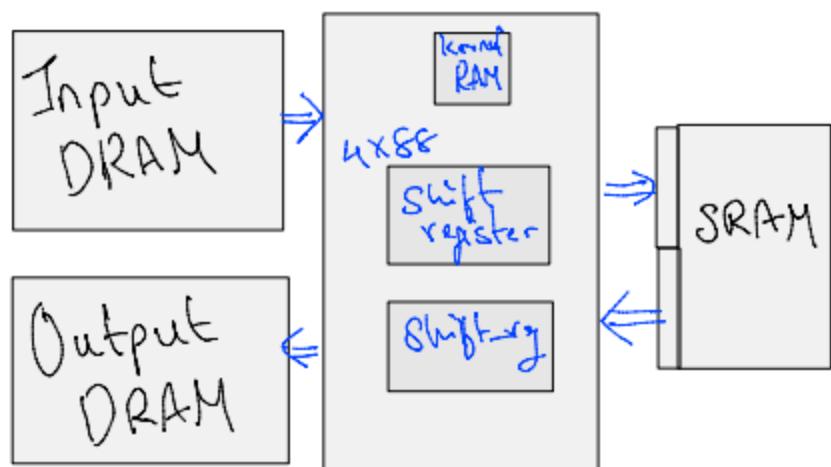
4. DRAM Output Writer FSM

Packs 8 processed output pixels into a burst and writes to memory with correct padding.

3. Design Architecture

3.1 Top-Level Block Diagram

Top level Design

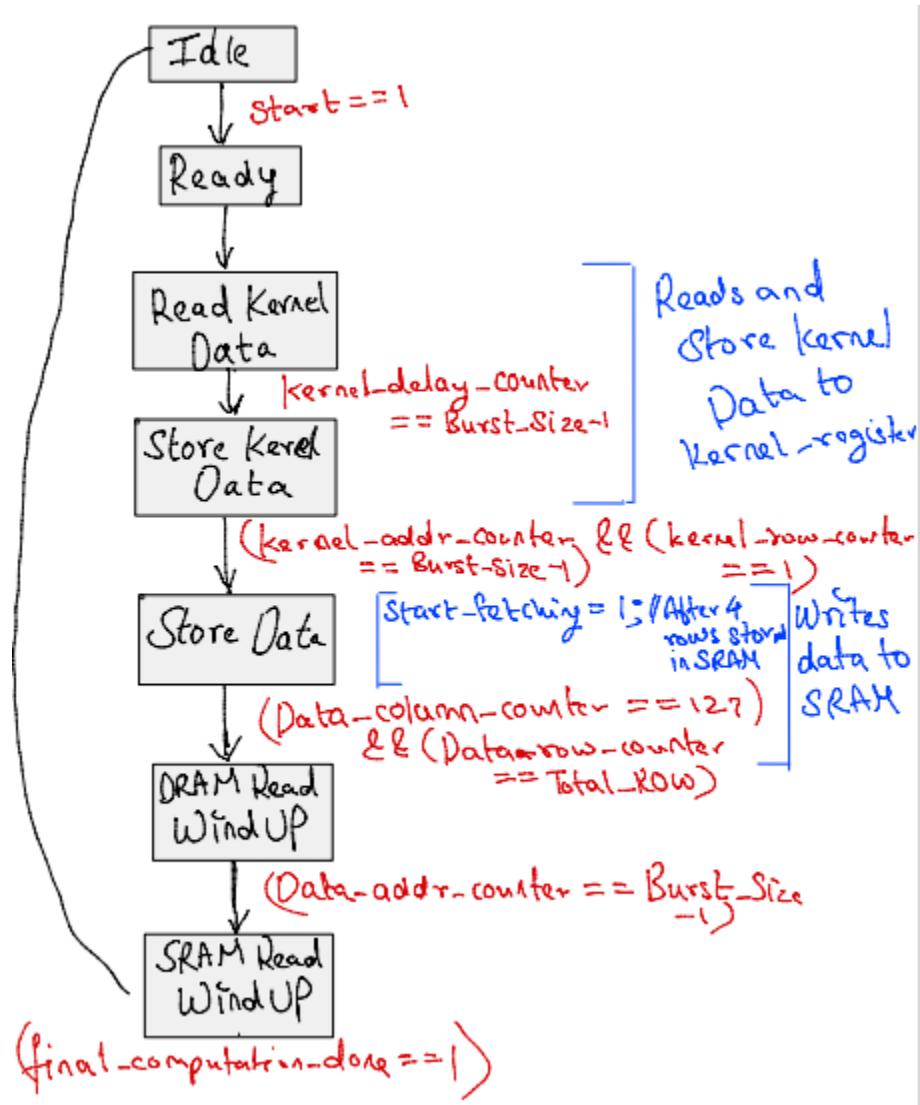


Note: Shift registers are used to compute the convolution of two rows at a time

4. Controller Architecture (FSMs)

I use four interacting FSMs, each corresponding to a different phase.

4.1 Input FSM (i_state)



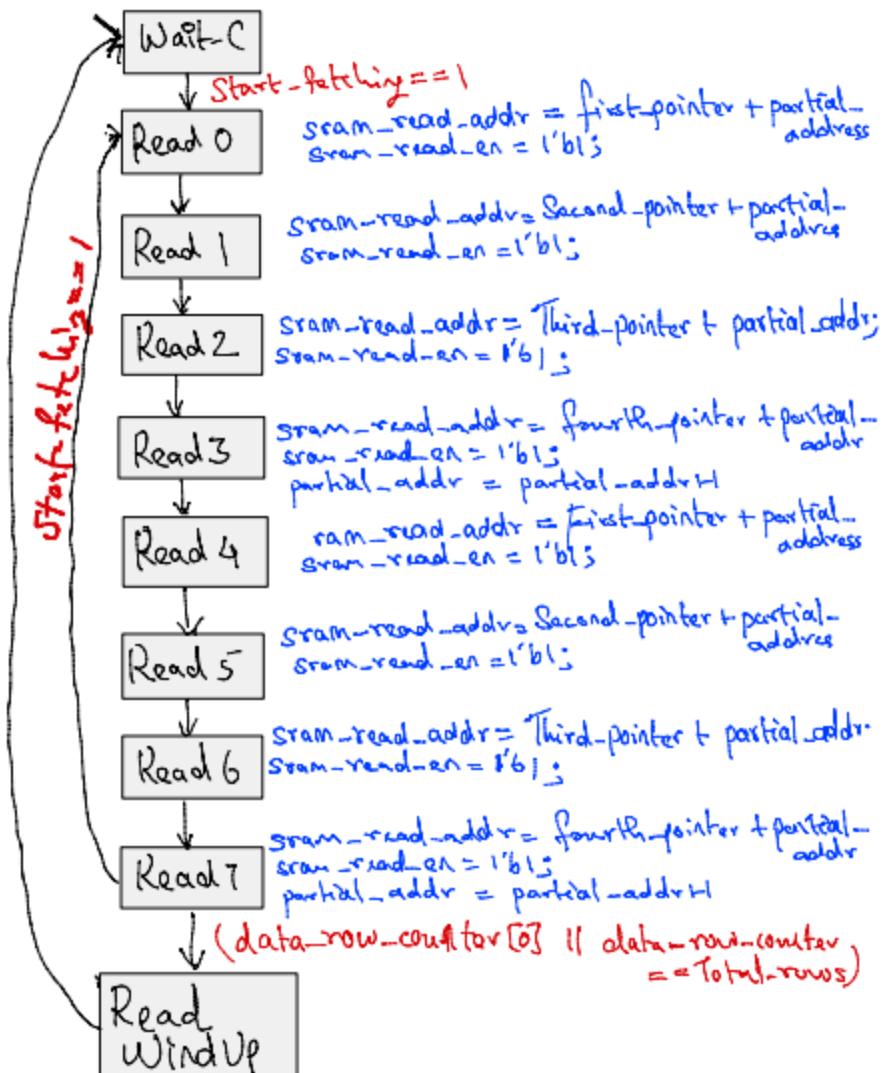
Controls DRAM reads:

State	Description
IDLE	ready=1, waiting for start

READY	Issues first kernel read
KERNEL_READ_DATA	Fetch 4×4 kernel
KERNEL_STORE_DATA	Write kernel bytes into kernel RAM
DATA_STORE	Read all image rows; write into SRAM
DRAM_READ_WINDUP	Flush final row of input reads
SRAM_READ_WINDUP	Wait for compute & output to finish

4.2 SRAM Read FSM (c_state)

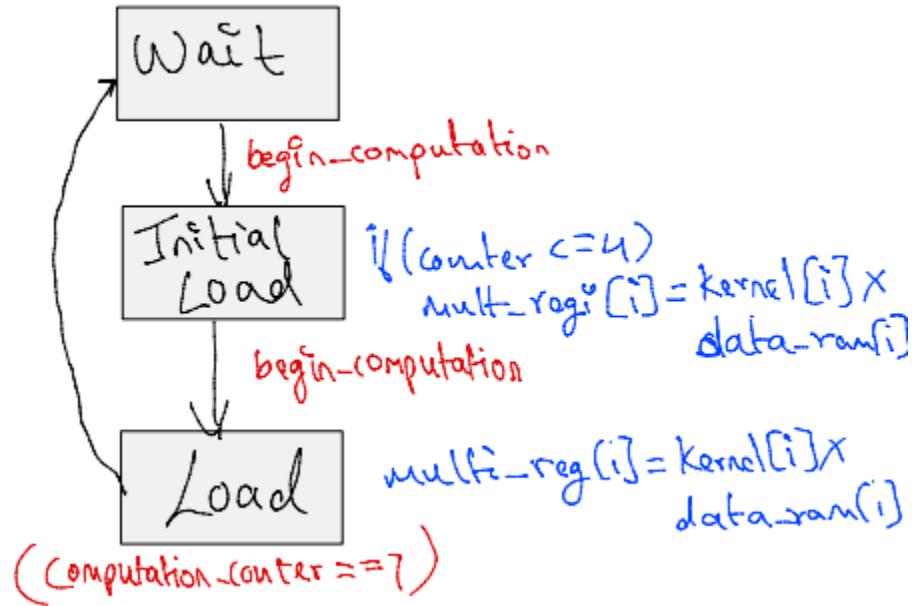
Cycles through 8 SRAM reads to fetch two consecutive 4-pixel segments for convolution.



State	Description
WAIT	idle
SRAM_READ0-7	fetch 8 required segments per window
WINDUP_READ	stall alignment
LAST_READ_WINDUP	final cleanup

The FSM ensures continuous data delivery to the compute engine.

4.3 Convolution/Pooling FSM (m_state)

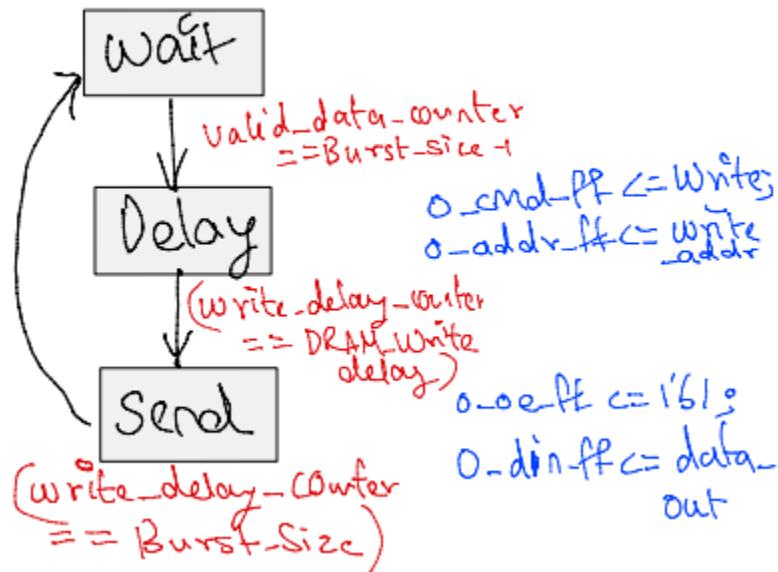


State	Description
WAIT_M	idle until read buffer is ready
INITIAL_LOAD	first compute cycle for new row
LOAD0	pipeline processing of sliding windows

Architecture uses:

- `multi_result1/2[0:15]` (parallel multipliers)
- 3-level adder tree (`partial_sum1/2/3`)
- Post-adder adjustment for ReLU
- Shift-based division for pooling

4.4 Output FSM (o_state)



State	Description
WAIT_O	waiting for 8 pooled output values
WAIT_DELAY	DRAM write delay alignment
SEND_DATA	sequentially send 8 bytes to DRAM

This FSM ensures correct padding and write-back alignment.

5. Datapath Description

5.1 Kernel Storage

The kernel is read during the first 16 DRAM bytes and stored into:

```
logic signed [7:0] kernel_ram [0:15];
```

Addressing matches the fixed pattern described in the spec.

5.2 SRAM Buffer Organization

SRAM stores 1024-byte rows in two 32-bit words:

```
Row i → SRAM entries:  
[i*256 + col] contains 4 consecutive pixels
```

Two buffers (`main_sram_buff_0` and `_1`) hold sliding window data.

5.3 Convolution Engine

Parallel Multiply Stage

16 independent multipliers:

```
multi_result1[i] = kernel_ram[i] * source_data_1[i];  
multi_result2[i] = kernel_ram[i] * source_data_2[i];
```

Adder Tree (Pipelined)

The design performs:

1. Level-1: 8 adds
2. Level-2: 4 adds
3. Level-3: 2 adds
4. Final sum: 1 add

This produces **20-bit convolution output**.

5.4 Leaky-ReLU

Specification:

- $x > 0 \rightarrow x$

- $-4 < x \leq 0 \rightarrow 0$
- $x \leq -4 \rightarrow$ arithmetic shift right by 2

Exactly implemented in:

```
if(convolution_result1_temp > 0) ...
else if(convolution_result1_temp <= 0 && convolution_result1_temp > -4)
    convolution_result1 <= 0;
else convolution_result1 <= ((convolution_result1_temp+3) >>> 2);
```

5.5 2x2 Average Pooling

Pooling sums two convolution outputs across two rows:

```
partial_avrg accumulates 2x2 region
final_avrg <= partial_avrg >>> 2; // custom truncation rule
```

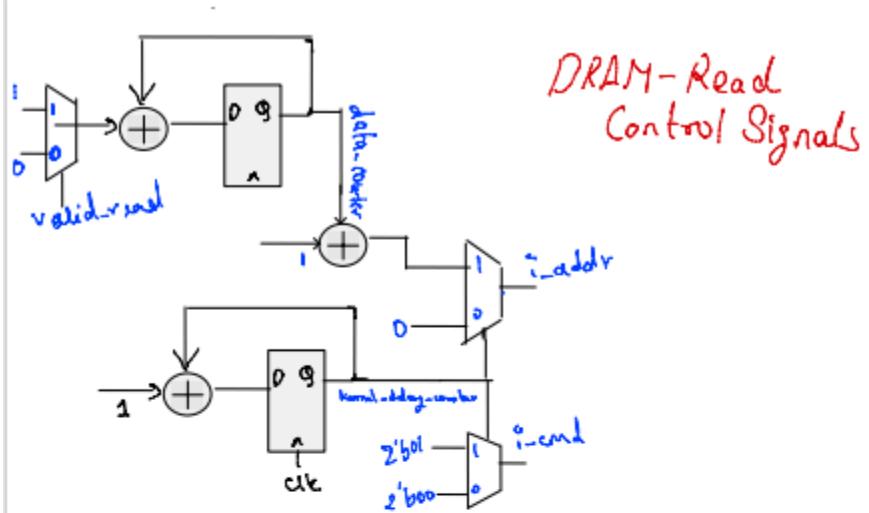
This matches the project's modified integer-pooling definition.

5.6 Output Clamping

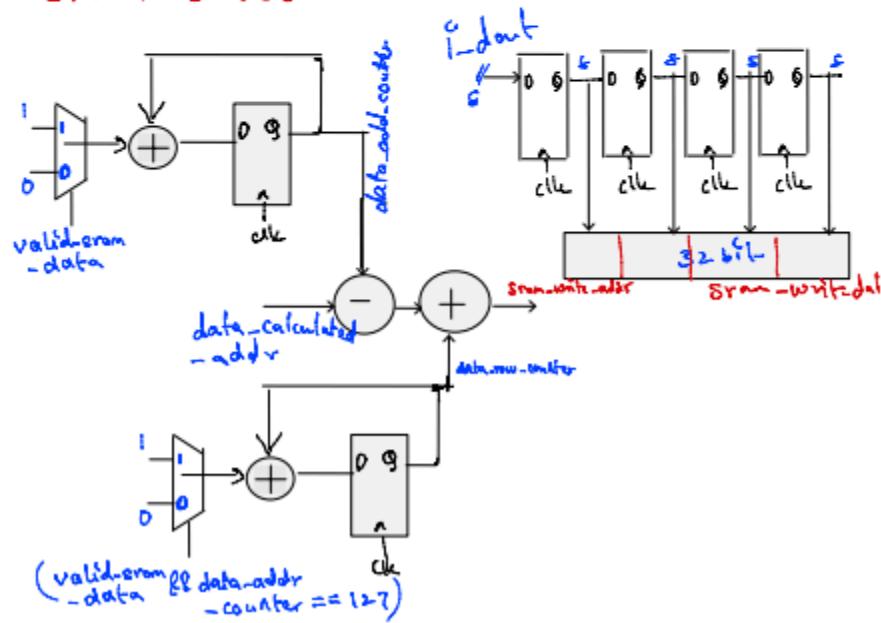
Final values limited to 8-bit signed:

```
> +127 → 127
< -128 → -128
else: keep value
```

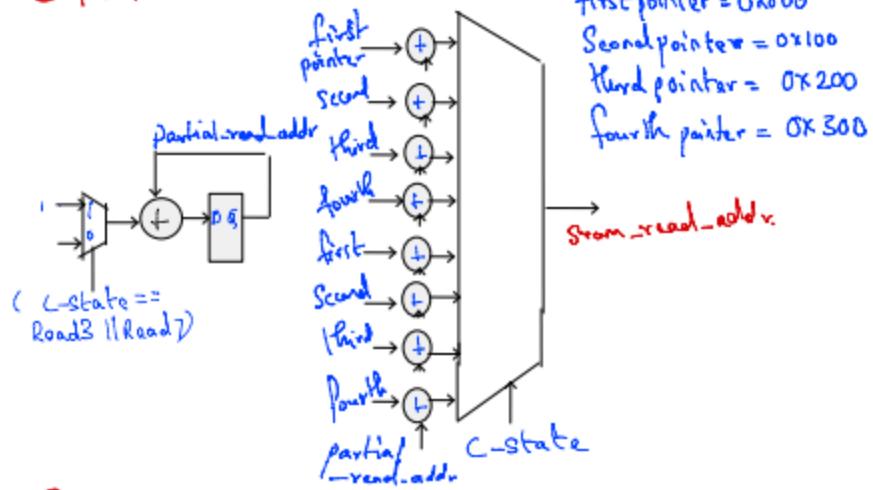
6. Schematics



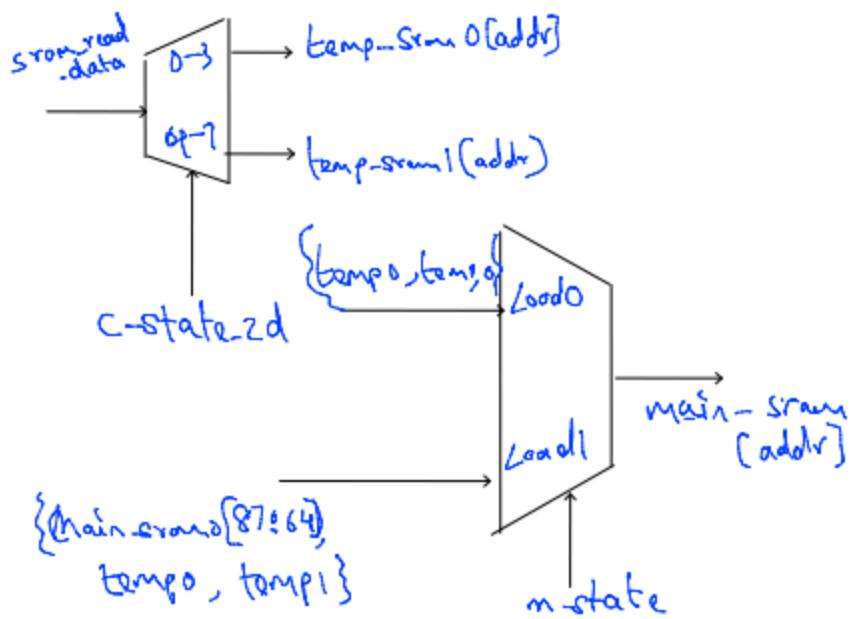
SRAM Write:

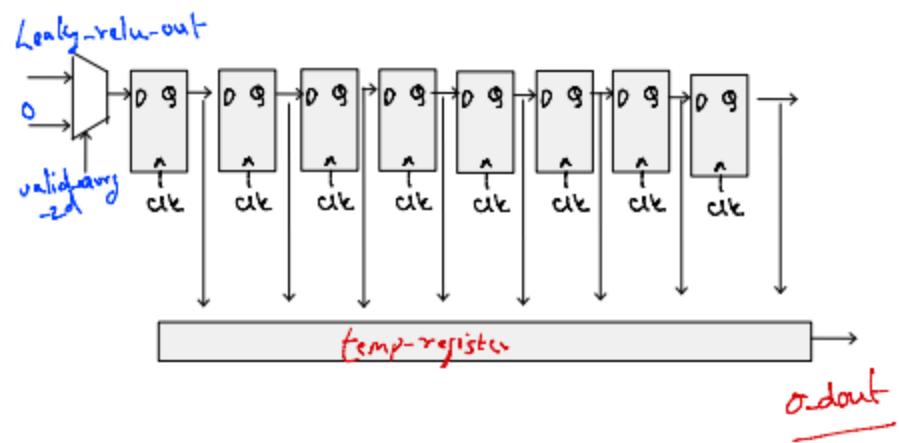
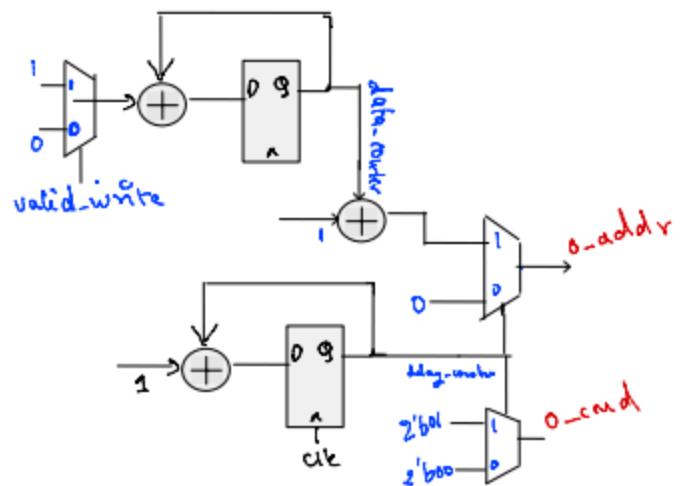


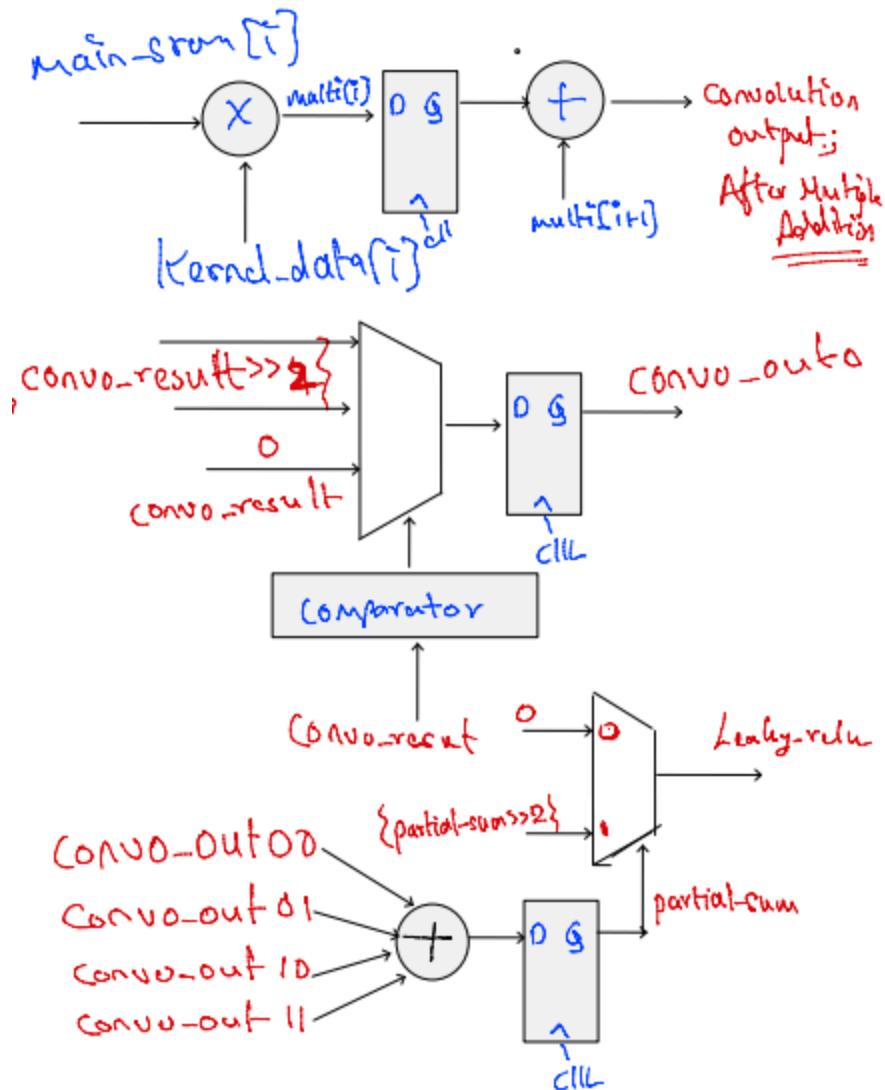
SRAM Read



Computation







7. Performance Characteristics

Metric	Value
Clock period	3.8 ns
Total Simulation cycles	12584124 cycles (According to printed results from modelsim) Actual for one run : 1,048,677

Cycles for One testcase	1,048,677 cycles or (2097354 as per results from Modelsim)
Cell area	37905
Worst-case setup slack	0.0001
Worst-case hold slack	0.0217

cell_report_final.rpt timing_max_slow_holdfixed_tut1.rpt

```
# A fanout number of 1000 was used for high fanout net computations.

Operating Conditions: slow   Library: NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm
Wire Load Model Mode: top

Startpoint: kernel_ram_reg[4][5]
            (rising edge-triggered flip-flop clocked by clk)
Endpoint: f_multi_result2_reg[4][13]
            (rising edge-triggered flip-flop clocked by clk)
Path Group: clk
Path Type: max

Point                                Incr      Path
-----
clock clk (rise edge)                0.0000    0.0000
clock network delay (ideal)          0.0000    0.0000
kernel_ram_reg[4][5]/CK (DFF_X1)     0.0000 #  0.0000 r
kernel_ram_reg[4][5]/Q (DFF_X1)      0.4904    0.4904 r
U22056/Z (BUF_X4)                  0.3671    0.8575 r
U20430/Z (XOR2_X2)                 0.4199    1.2774 r
U21609/ZN (NAND2_X2)               0.0768    1.3542 f
U20788/ZN (NAND2_X2)               0.1713    1.5255 r
U21608/ZN (INV_X2)                 0.0514    1.5769 f
U21084/ZN (NAND2_X2)               0.0721    1.6490 r
U6476/ZN (NAND2_X1)                0.0720    1.7210 f
U20960/ZN (NAND2_X2)               0.1133    1.8344 r
U8501/ZN (XNOR2_X1)                0.3524    2.1868 r
U10742/ZN (XNOR2_X2)               0.3342    2.5210 r
U20344/ZN (NAND2_X1)               0.1222    2.6432 f
U13431/ZN (INV_X1)                 0.1190    2.7622 r
U18055/ZN (NAND2_X1)               0.0749    2.8371 f
U10836/ZN (NAND3_X2)               0.1901    3.0272 r
U6782/ZN (AOI21_X1)                0.1120    3.1392 f
U10835/ZN (NAND3_X2)               0.2402    3.3794 r
U10834/ZN (NAND2_X2)               0.0679    3.4473 f
f_multi_result2_reg[4][13]/D (DFF_X2) 0.0000    3.4473 f
data arrival time                   3.4473

clock clk (rise edge)                3.8000    3.8000
clock network delay (ideal)          0.0000    3.8000
clock uncertainty                   -0.0500    3.7500
f_multi_result2_reg[4][13]/CK (DFF_X2) 0.0000    3.7500 r
library setup time                  -0.3026    3.4474
data required time                  3.4474

data required time                  3.4474
data arrival time                   -3.4473

slack (MET)                         0.0001
```

cell_report_final.rpt	timing_max_slow_holdfixed_tut1.rpt
DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_computation_pipelined_reg[2] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_computation_pipelined_reg[3] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_computation_pipelined_reg[4] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_computation_pipelined_reg[5] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_computation_pipelined_reg[6] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_data_counter_d_reg[0] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_data_counter_d_reg[1] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_data_counter_d_reg[2] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_data_counter_reg[0] DFF_X2	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_data_counter_reg[1] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_data_counter_reg[2] DFF_X2	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_data_d_reg DFF_X2	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_kernel_data_d_reg DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_sram_read_2d_reg DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
valid_sram_read_d_reg DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
write_delay_counter_reg[0] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
write_delay_counter_reg[1] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
write_delay_counter_reg[2] DFF_X1	NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm 4.7880 n
Total 27446 cells	37905.7981
1	

```

cell_report_final.rpt          timing_max_slow_holdfixed_tut1.rpt      timing_max_slow.rpt
# A fanout number of 1000 was used for high fanout net computations.

Operating Conditions: slow Library: NangateOpenCellLibrary_PDKv1_2_v2008_10_slow_nldm
Wire Load Model Mode: top

Startpoint: kernel_ram reg[4][5]
(rising edge-triggered flip-flop clocked by clk)
Endpoint: f_multi_result2_reg[4][13]
(rising edge-triggered flip-flop clocked by clk)
Path Group: clk
Path Type: max

Point                                     Incr     Path
-----
clock clk (rise edge)                   0.0000   0.0000
clock network delay (ideal)           0.0000   0.0000
kernel_ram reg[4][5]/CK (DFF_X1)       0.0000 # 0.0000 r
kernel_ram reg[4][5]/Q (DFF_X1)       0.4904   0.4904 r
U22056/Z (BUF_X4)                     0.3671   0.8575 r
U20438/Z (XOR2_X2)                   0.4199   1.2774 r
U21609/ZN (NAND2_X2)                 0.0768   1.3542 f
U20788/ZN (NAND2_X2)                 0.1713   1.5255 r
U21608/ZN (INV_X2)                   0.0514   1.5769 f
U21084/ZN (NAND2_X2)                 0.0721   1.6490 r
U6476/ZN (NAND2_X1)                  0.0720   1.7210 f
U20966/ZN (NAND2_X2)                 0.1133   1.8344 r
U8501/ZN (XNOR2_X1)                  0.3524   2.1868 r
U10742/ZN (XNOR2_X2)                 0.3342   2.5210 r
U20344/ZN (NAND2_X1)                 0.1222   2.6432 f
U13431/ZN (INV_X1)                   0.1190   2.7622 r
U18055/ZN (NAND2_X1)                 0.0749   2.8371 f
U10836/ZN (NAND3_X2)                 0.1901   3.0272 r
U6782/ZN (A0121_X1)                  0.1120   3.1392 f
U10835/ZN (NAND3_X2)                 0.2402   3.3794 r
U10834/ZN (NAND2_X2)                 0.0679   3.4473 f
f_multi_result2_reg[4][13]/D (DFF_X2) 0.0000   3.4473 f
data arrival time                      0.0000   3.4473

clock clk (rise edge)                   3.8000   3.8000
clock network delay (ideal)           0.0000   3.8000
clock uncertainty                     -0.0500   3.7500
f_multi_result2_reg[4][13]/CK (DFF_X2) 0.0000   3.7500 r
library setup time                    -0.3026   3.4474
data required time                   3.4474
data arrival time                      3.4473

slack (MET)                           0.0001

```

Documents Open timing_min_fast_holdcheck_tut1.rpt

cell_report_final.rpt x cell_report_final.rpt x timing_max_slow_holdfixed_tut1.rpt x timing_max_slow.rpt x timing_min_fast_holdcheck_tut1.rpt x

Information: Updating design information... (UID-85)
Warning: Design 'dut' contains 1 high-fanout nets. A fanout number of 1000 will be used for delay calculations involving these nets. (TIM-134)

Report : timing
-path full
-delay min
-max_paths 1
Design : dut
Version: T-2022-03-SP4
Date : Tue Nov 18 11:25:58 2025

A fanout number of 1000 was used for high fanout net computations.

Operating Conditions: fast Library: NangateOpenCellLibrary_PDKv1_2_v2008_10_fast_nldm
Wire Load Model Mode: top

Startpoint: temp_sram_buff0_1_reg[3][8]
(rising edge-triggered flip-flop clocked by clk)
Endpoint: main_sram_buff_0_reg[3][29]
(rising edge-triggered flip-flop clocked by clk)
Path Group: clk
Path Type: min

Point Incr Path

clock clk (rise edge) 0.0000 0.0000
clock network delay (ideal) 0.0000 0.0000
temp_sram_buff0_1_reg[3][5]/CK (DFF_X2) 0.0000 # 0.0000 r
temp_sram_buff0_1_reg[3][5]/Q (DFF_X2) 0.0541 0.0541 r
U20870/ZN (A0121_X1) 0.0176 0.8716 f
main_sram_buff_0_reg[3][29]/D (DFF_X1) 0.0000 0.8716 f
data arrival time 0.0000 0.8716

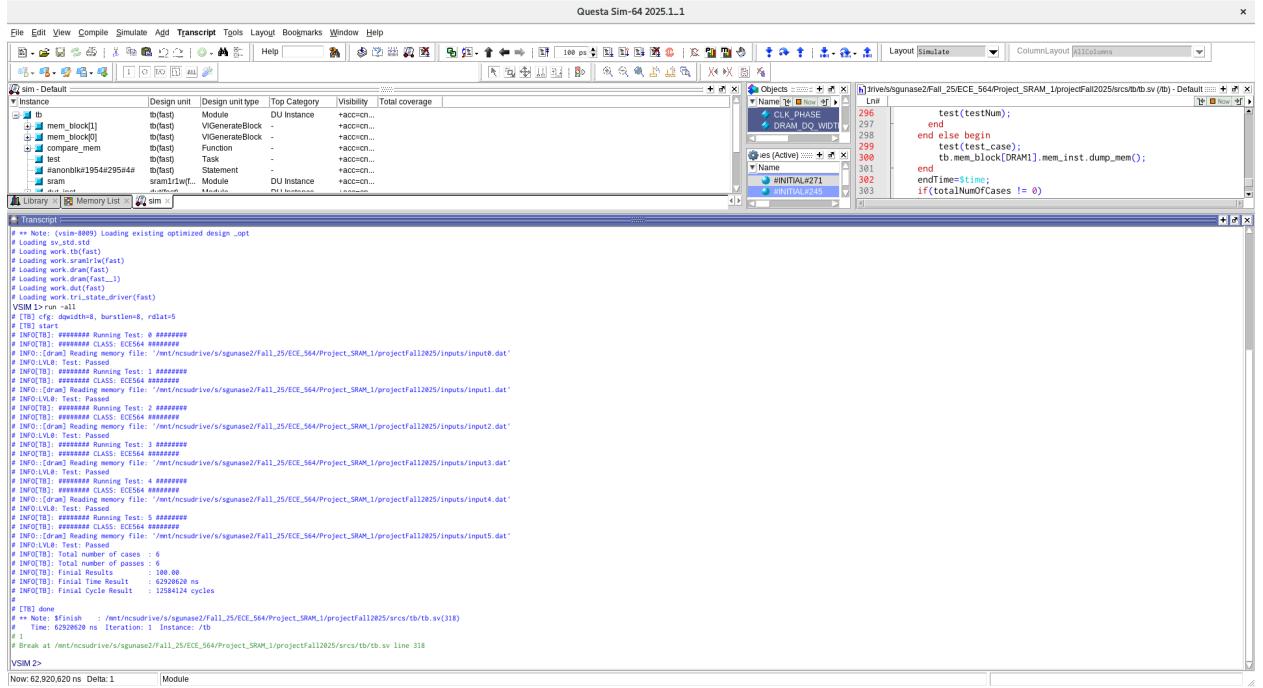
clock clk (rise edge) 0.0000 0.0000
clock network delay (ideal) 0.0000 0.0000
clock uncertainty 0.0500 0.0500
main_sram_buff_0_reg[3][29]/CK (DFF_X1) 0.0000 0.0500 r
library hold time -0.0001 0.0499
data required time 0.0499
data arrival time 0.0499

slack (MET) 0.0217

Plain Text Tab Width: 8 Ln 1, Col 1 INS

8. Correctness Verification

All six test cases passed and there were no major synthesis error.



9. Conclusion

This project implements a fully pipelined CNN hardware accelerator for a 1024×1024 image using:

- 4×4 convolution
 - Leaky-ReLU
 - 2×2 average pooling
 - Fully streaming pipeline aligned to DRAM bandwidth
 - SRAM-assisted sliding window engine
 - 7-stage computation pipeline

The design meets all memory/timing specifications and utilizes SystemVerilog features such as `typedef enum`, `always_ff`, and multi-dimensional packed arrays.